Autoscaling Configuration for Backend and Frontend Infrastructure on AWS

Backend Autoscaling

Step 1: Setting up the Backend AMI

1. **Launch a new EC2 instance**:
    - Install and configure the backend application (Node.js, Django, etc.) and any dependencies.
    - Validate the setup to ensure the backend functions correctly on HTTP port 8000.
2. **Create an AMI from the EC2 instance**:
    - Go to **EC2 Dashboard** > **Instances**.
    - Select the instance and click **Actions** > **Create Image**.
    - Name the AMI appropriately (e.g., Backend_AMI).
    - This AMI will serve as the template for launching instances in the Auto Scaling group.

Step 2: Create Launch Template for Backend Instances

1. **Go to Launch Templates** in the **EC2 Dashboard**.
2. **Create Launch Template**:
    - Provide a name (e.g., Backend_Launch_Template).
    - Under **AMI ID**, select the Backend_AMI created in Step 1.
    - Choose an instance type appropriate for the backend.
    - Configure the template to use the backend security group (created in Step 5).
    - Optionally, configure other details like network settings and IAM roles.

Step 3: Setting up the Target Group for Backend Instances

1. **Create a Target Group**:
    - Go to **Target Groups** in the **EC2 Dashboard** > **Create Target Group**.
    - Choose **Instances** as the target type.

- Set **Protocol** to **HTTP** and **Port** to **8000**.
- Select the appropriate **VPC** and configure **Health Checks** to monitor the health of instances on HTTP port 8000.
- Leave the **Register Targets** section blank, as instances will be added dynamically by the Auto Scaling Group.

Step 4: Configuring the Internal Application Load Balancer for Backend

1. **Create an Internal ALB**:
   - Go to **Load Balancers** > **Create Load Balancer**.
   - Select **Application Load Balancer** and set it to **Internal** for VPC-only access.
   - Configure listeners to listen on **Port 80**.
2. **Attach the Backend Target Group to the ALB**:
   - Under **Listeners and Routing**, add a rule to forward traffic from **Port 80** to the **Backend Target Group** created in Step 3.
   - Assign the ALB to **private subnets** across multiple availability zones for high availability.

Step 5: Configure Backend Security Group

1. **Create a Security Group for the Backend ALB**:
   - Go to **Security Groups** in **VPC Dashboard**.
   - Create a group (e.g., backend-alb-sg) that allows **HTTP (port 80)** traffic from the frontend ALB security group.
2. **Configure Security Group for Backend EC2 Instances**:
   - Create another Security Group for backend instances (e.g., backend-instances-sg).
   - Allow **HTTP (port 8000)** traffic from the backend ALB security group.
   - Restrict outbound access as necessary to maintain security.

Step 6: Create an Autoscaling Group for the Backend

- Go to **EC2 Console** > **Auto Scaling Groups** > **Create Auto Scaling Group**.
- Use the launch template created in Step 2.
- Attach the internal load balancer and backend target group to ensure load distribution.
- Set desired capacity, scaling policies, and health checks as needed.
- Complete the configuration to launch and manage backend instances.

Frontend Autoscaling

Step 1: Setting up the Frontend AMI

1. **Launch a new EC2 instance**:
    - Set up the frontend application (Apache, NGINX, etc.) and any required configurations.
    - Validate that the frontend application works correctly on HTTP port 80.
2. **Create an AMI from the EC2 instance**:
    - Go to **EC2 Dashboard** > **Instances**.
    - Select the instance and click **Actions** > **Create Image**.
    - Name this AMI (e.g., Frontend_AMI), which will serve as the template for the frontend Auto Scaling Group.

Step 2: Create Launch Template for Frontend Instances

1. **Go to Launch Templates** in the **EC2 Dashboard**.
2. **Create Launch Template**:
    - Provide a name (e.g., Frontend_Launch_Template).
    - Select the Frontend_AMI created in Step 1.
    - Choose an instance type suitable for the frontend.
    - Configure the template to use the frontend security group (configured in Step 5).

Step 3: Setting up the Target Group for Frontend Instances

1. **Create a Target Group for Frontend**:
    - Go to **Target Groups** in **EC2 Dashboard** > **Create Target Group**.
    - Select **Instances** as the target type and set **Protocol** to **HTTP** with **Port 80**.
    - Configure **Health Checks** to monitor HTTP traffic on Port 80.
    - Leave **Register Targets** empty to let the Auto Scaling Group handle instance registration.

Step 4: Configuring the Public Application Load Balancer for Frontend

1. **Create a Public ALB**:
    - Go to **Load Balancers** > **Create Load Balancer**.
    - Choose **Application Load Balancer** and select **Internet-facing** to allow external access.
    - Configure listeners to listen on **Port 80**.
2. **Attach the Frontend Target Group to the Public ALB**:
    - Under **Listeners and Routing**, forward traffic from **Port 80** to the **Frontend Target Group** created in Step 3.
    - Assign the ALB to **public subnets** in multiple availability zones to ensure high availability.

Step 5: Configure Frontend Security Group

1. **Create a Security Group for the Public ALB**:
    - In **Security Groups** under **VPC Dashboard**, create a group (e.g., frontend-alb-sg).
    - Allow **HTTP (port 80)** from **0.0.0.0/0** to enable public internet access.
2. **Configure Security Group for Frontend EC2 Instances**:
    - Create another Security Group (e.g., frontend-instances-sg) for frontend instances.
    - Allow **HTTP (port 80)** traffic from the frontend ALB security group.
    - Limit outbound access as needed for security.

Step 6: Create an Autoscaling Group for the Frontend

- Go to **EC2 Console** > **Auto Scaling Groups** > **Create Auto Scaling Group**.
- Select the launch template created in Step 2 for frontend scaling.
- Attach the public load balancer and frontend target group.
- Set desired capacity, scaling policies, and health checks.
- Finalize settings to launch and manage frontend instances in the autoscaling group.

Verification and Testing

Step 1: Simulate High Load on an Instance

1. **SSH into an EC2 Instance in the Autoscaling Group:**
    - Use SSH to connect to one of the instances created by the autoscaling group.

      ssh -i /path/to/key.pem ec2-user@<instance_public_ip>

2. **Increase CPU Load with stress Command:**
    - Install stress if it's not already installed:

      sudo apt-get install -y stress # For Ubuntu
    - Run the following command to increase CPU load:

      stress --cpu 6 -t 600
    - This command will simulate high CPU usage across 6 cores for 10 minutes.
3. **Monitor CPU Utilization in Another Terminal:**
    - In a separate terminal, SSH into the same instance and run:

      top
    - Observe CPU utilization, and confirm that it crosses the 50% threshold, which should trigger the autoscaling policy.

Step 2: Verify Autoscaling Triggers New Instance Launches

1. **Check AWS Console for New Instances:**
    - Go to the **EC2 Console** > **Auto Scaling Groups** > **Activity** to monitor the scaling activities.
    - You should see new instances being launched by the autoscaling group due to high CPU usage on the original instance.

Step 3: Manually Test Autoscaling by Stopping the Original Instance

1. **Stop the Original Instance:**
    - In the **EC2 Console**, select the original instance, and stop it (do not terminate).
    - This will simulate an instance failure, causing the autoscaling group to replace the stopped instance.
2. **Check Application Accessibility:**
    - Access the application through the load balancer endpoint to ensure that the new instances created by autoscaling are handling the load.

- This confirms that the autoscaling configuration is correctly set to handle instances going down or high CPU loads.

Security group configuration

1.) backend instance (sg-05c996808c00f2fa9)

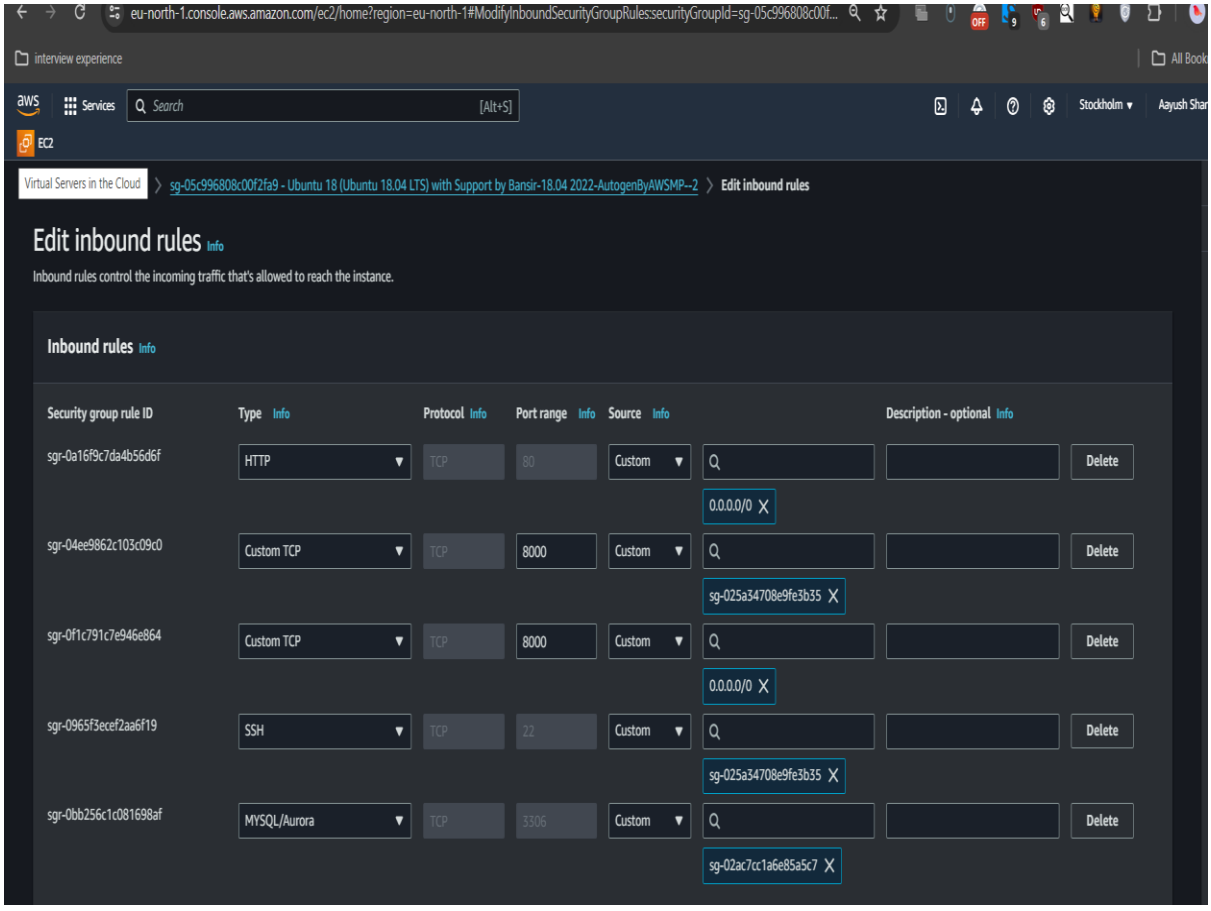

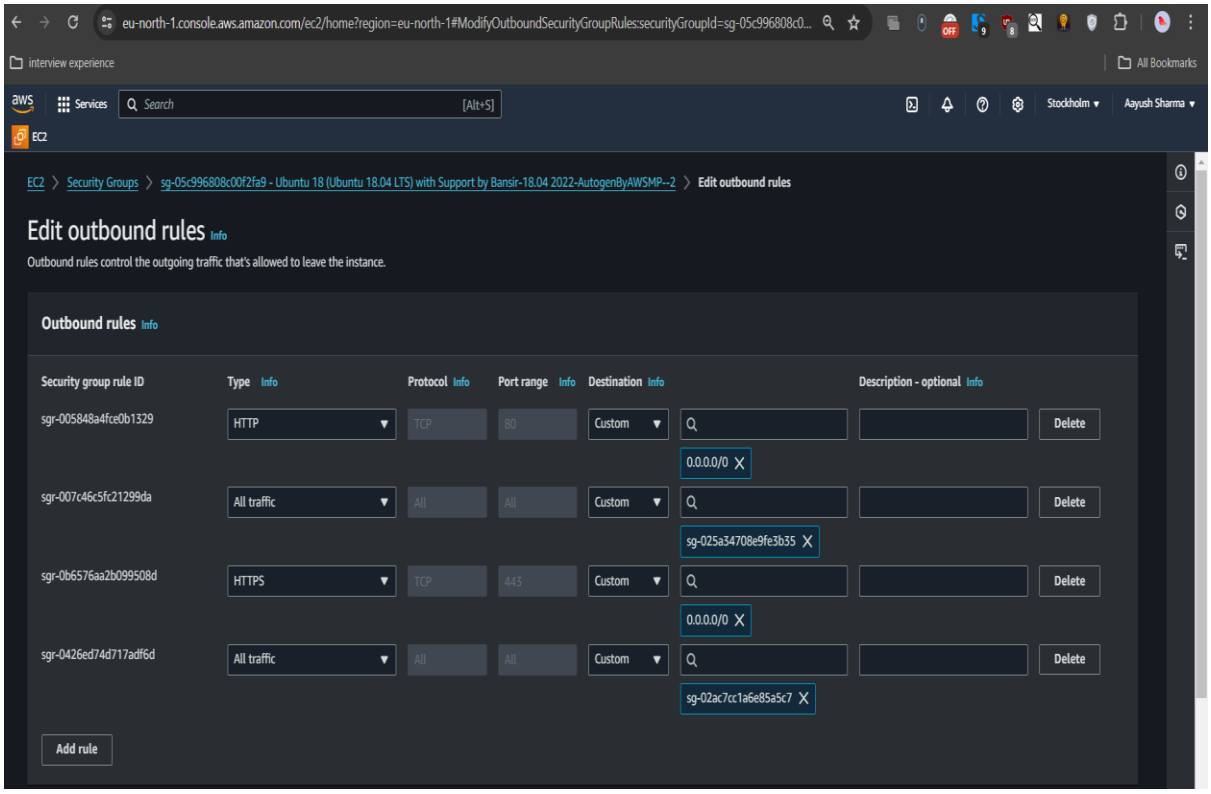

2.) Frontend instance (sg-025a34708e9fe3b35)

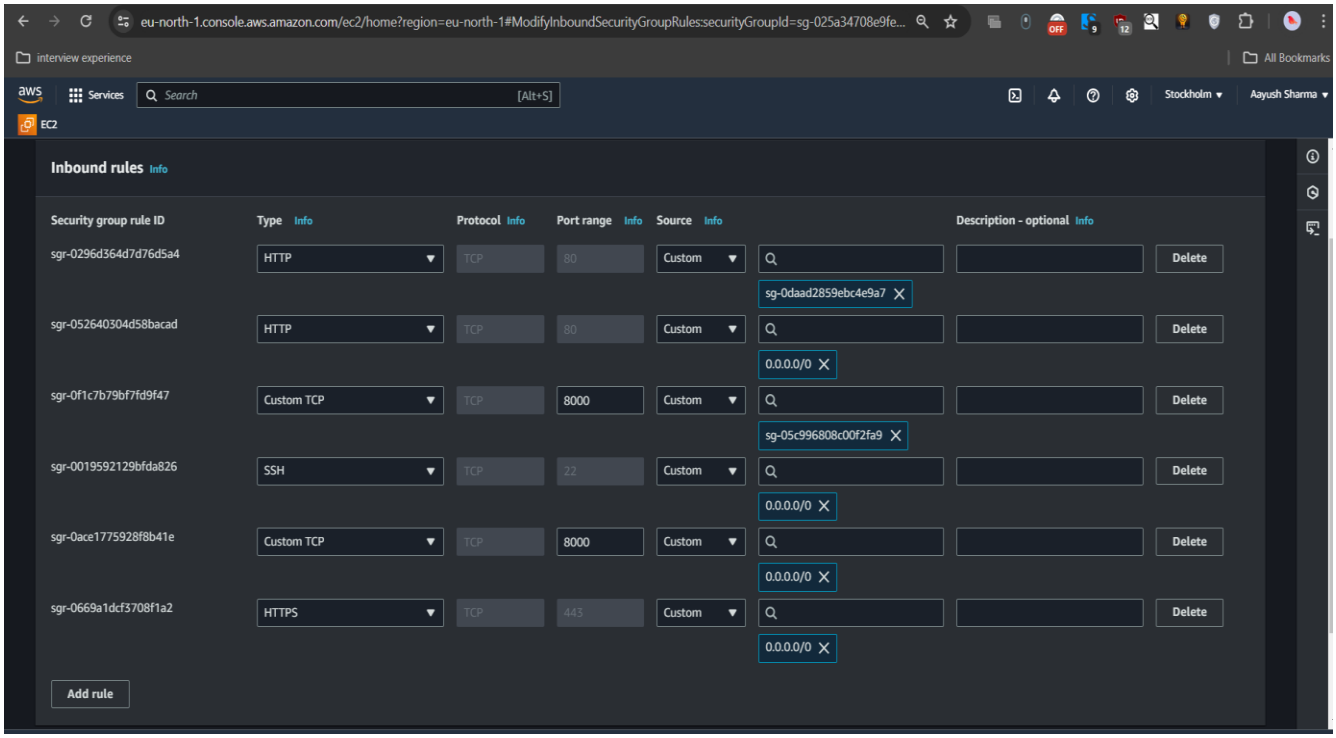

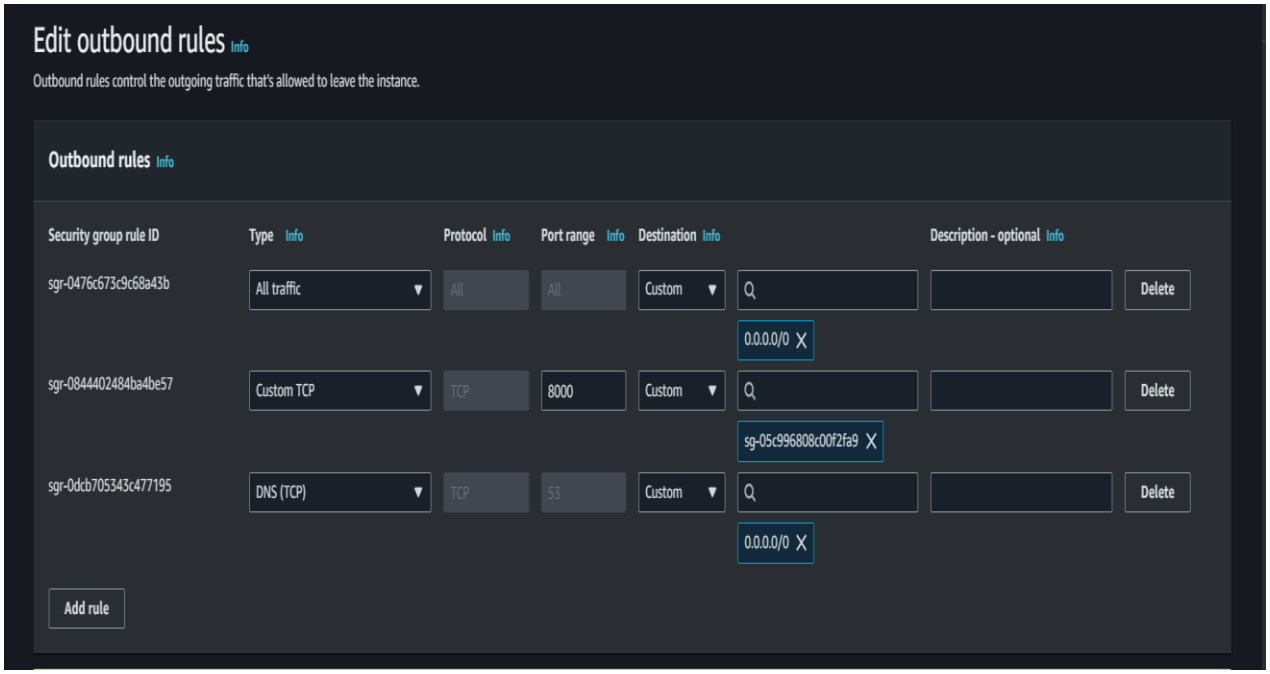

3.) database instance (sg-02ac7cc1a6e85a5c7)

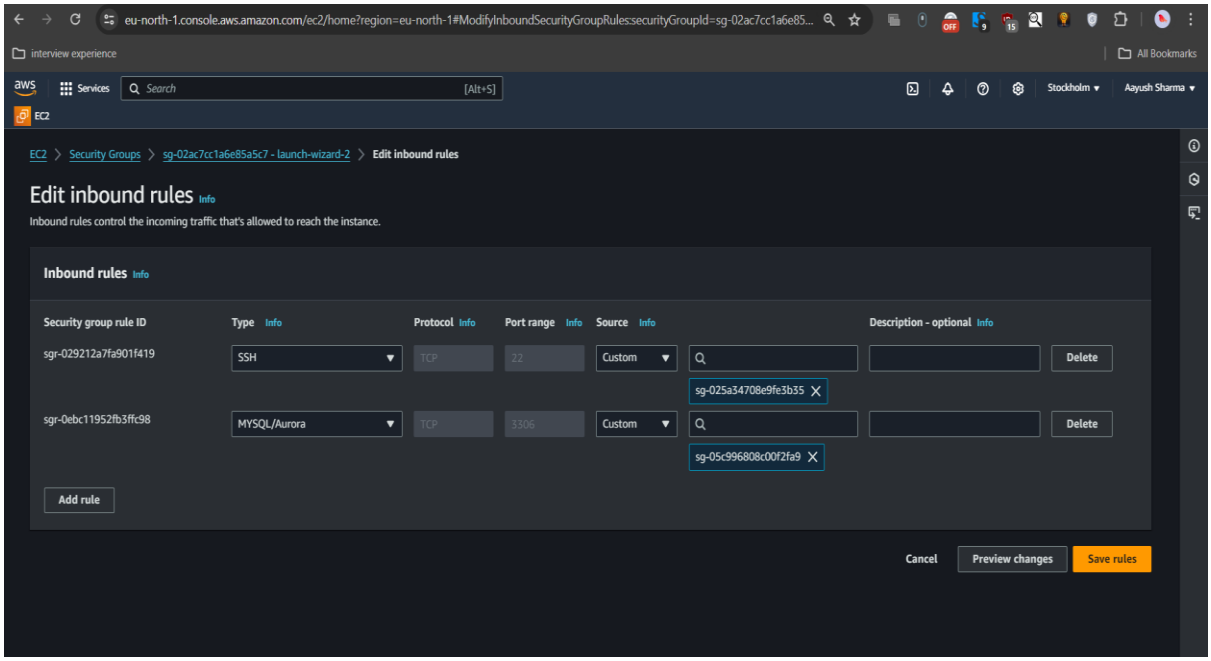

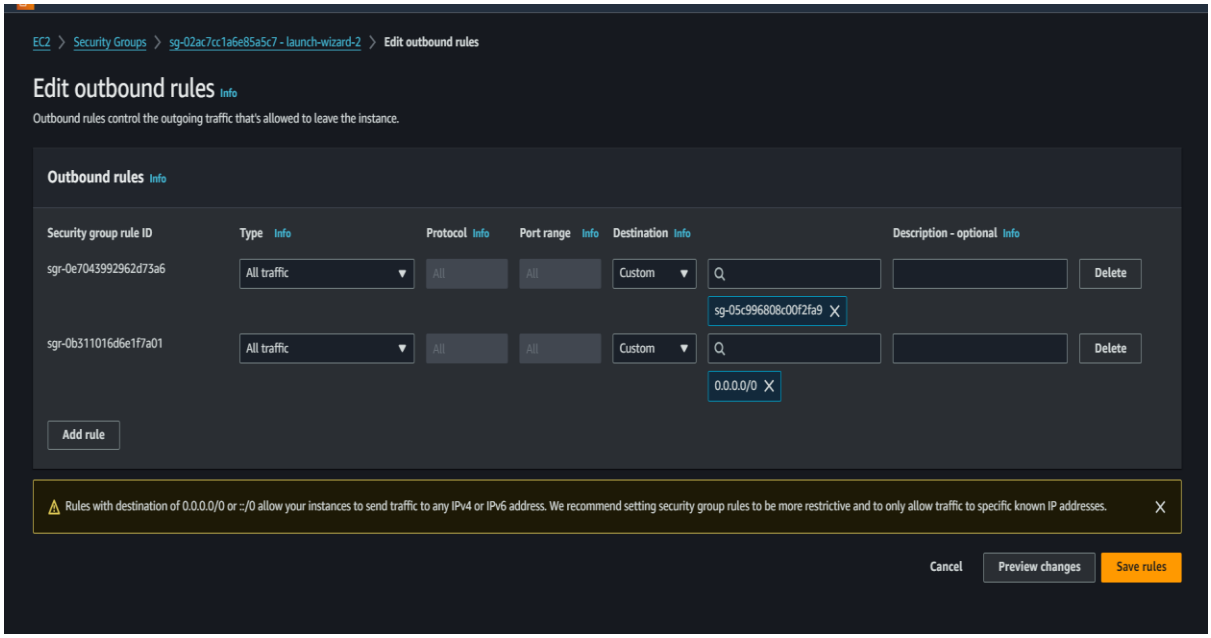

4.) backend loadbalancer (sg-0319f1c5b53b6d743)

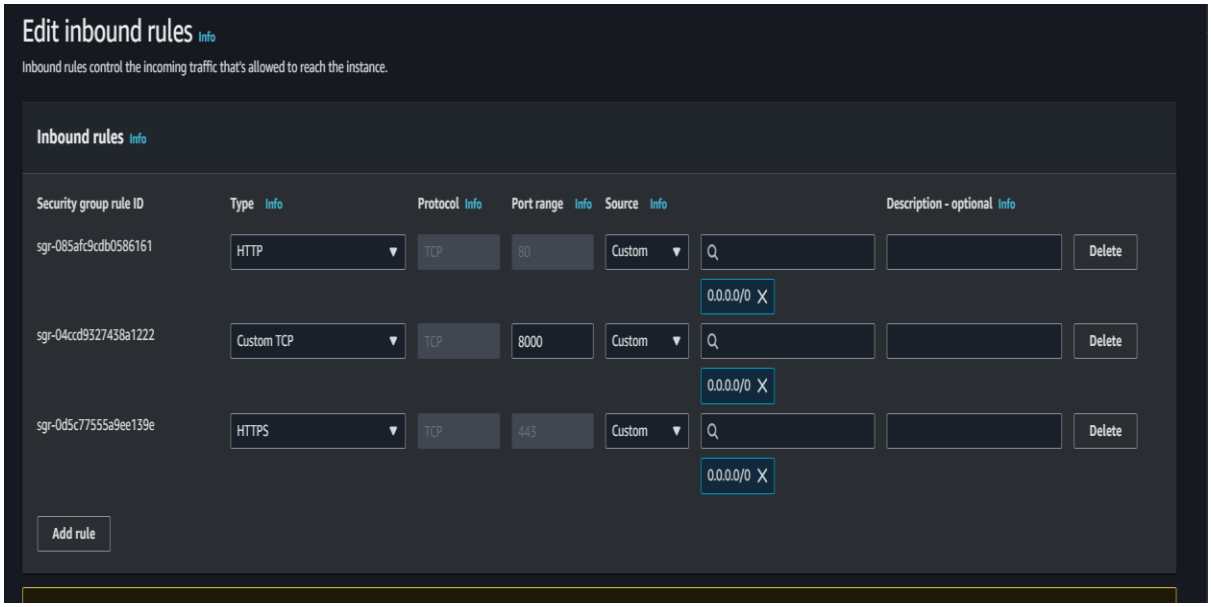

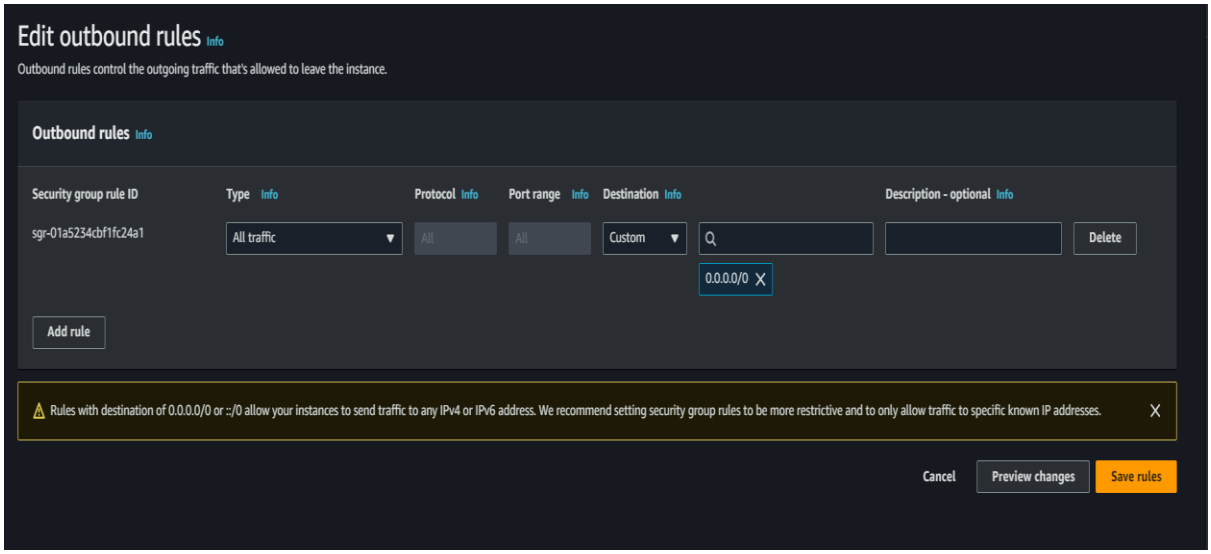

5.) frontend load balancer (sg-0daad2859ebc4e9a7)

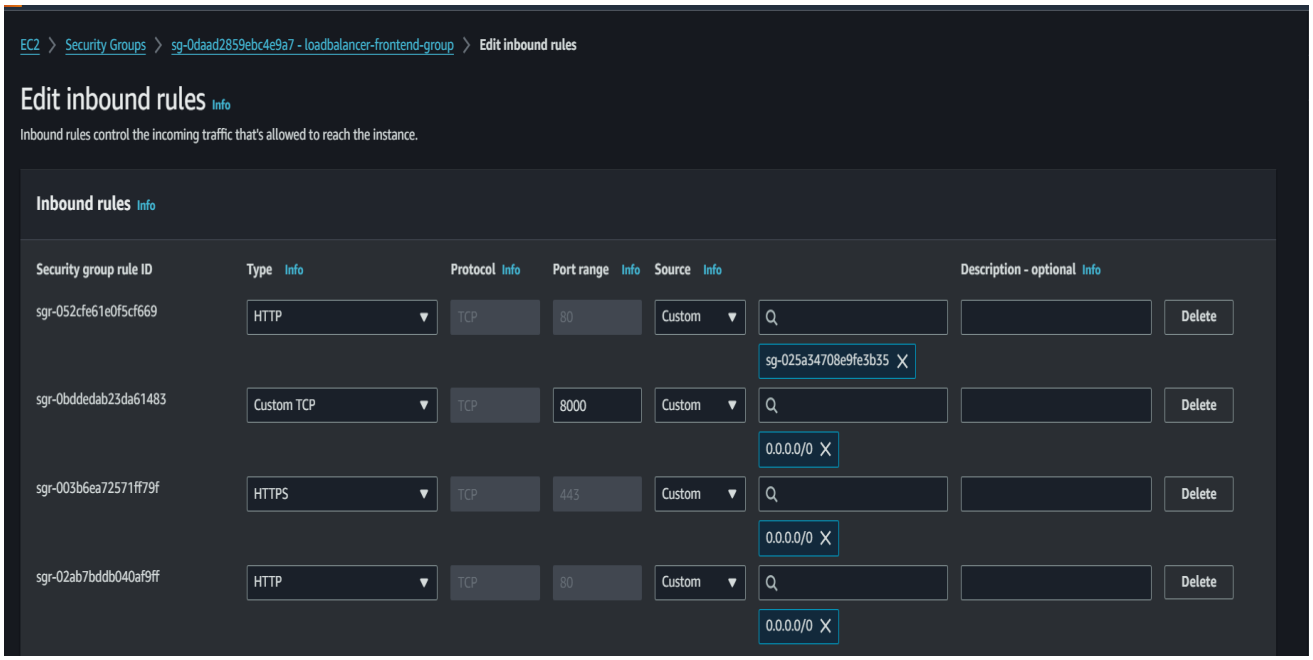

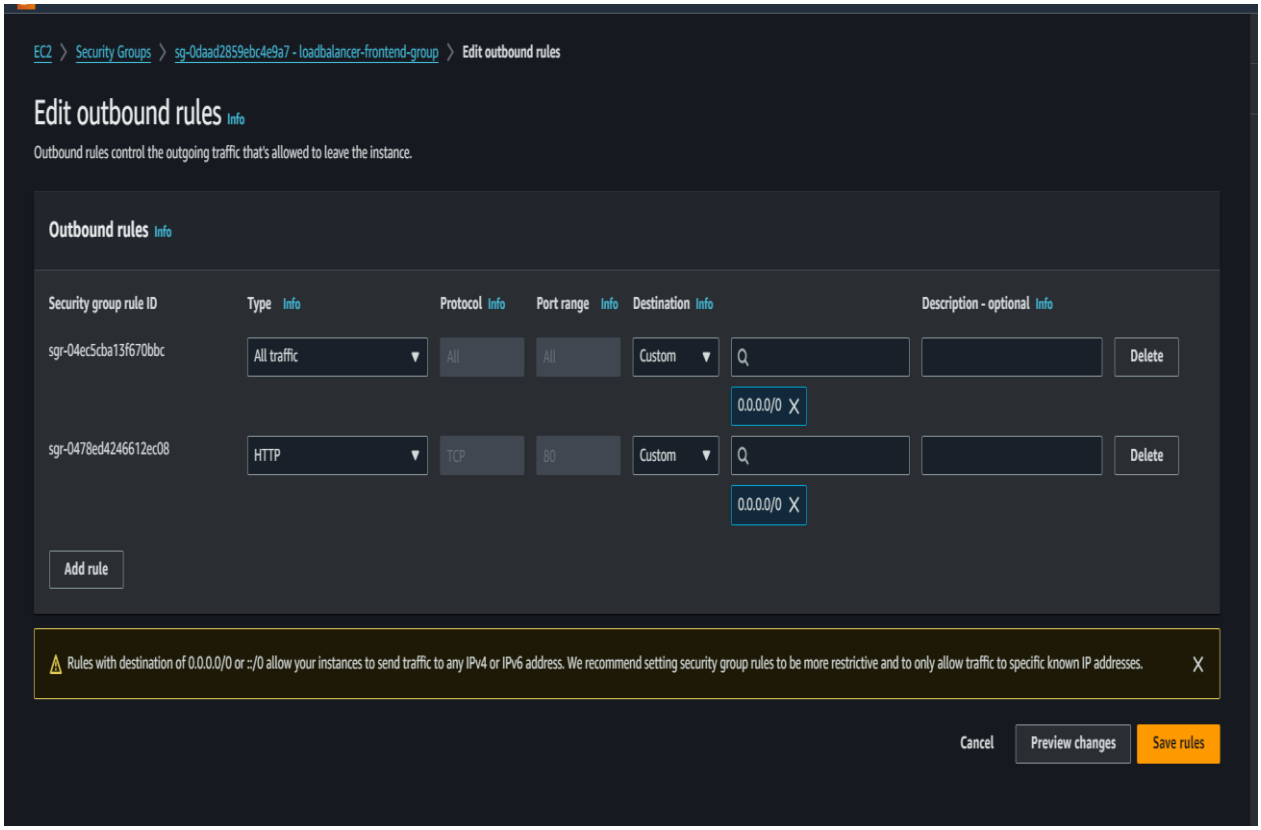

Nginx Configuration

```
upstream backend {
    server internal-backend-lb-354987202.eu-north-1.elb.amazonaws.com;  # AWS internal load balancer
    # Uncomment and add additional servers if necessary
    # server 10.0.2.88:8000;
}

server {
    listen 80;
    server_name frontend-lb-1116214595.eu-north-1.elb.amazonaws.com;

    # Serve static files
    location /static/ {
```

```
        alias /new_chatapp/fundoo/static/;  # Ensure this path is correct for static files
    }

    # Proxy pass backend requests to upstream block
    location / {
        proxy_pass http://backend;  # Reference the upstream block for backend requests
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
    }
}
```